

Q1) C

Q2) B

$$Q3) \quad p(Y=1|X) = \frac{e^{\theta_0 + \theta_1 X_1 + \theta_2 X_2}}{1 + e^{\theta_0 + \theta_1 X_1 + \theta_2 X_2}}$$

$$\hat{\theta}_0 = -6, \quad \hat{\theta}_1 = 0.05, \quad \hat{\theta}_2 = 1$$

$$\begin{aligned} a) \quad P(Y=1 | X_1=40, X_2=3.5) &= \frac{e^{-6+0.05(40)+1(3.5)}}{1 + e^{-6+0.05(40)+1(3.5)}} \\ &= 0.3775 \end{aligned}$$

$$b) \quad P(Y=1 | X_1, X_2=3.5) = 0.50$$

$$0.50 = \frac{e^{-6+0.05(X_1)+3.5}}{1 + e^{-6+0.05X_1+3.5}}$$

$$0.50 + 0.5e^{-2.5+0.05X_1} = e^{-2.5+0.05X_1}$$

$$0.5 = 0.5e^{-2.5+0.05X_1}$$

$$1 = e^{-2.5+0.05X_1}$$

$$0 = -2.5 + 0.05X_1$$

$$0.05X_1 = 2.5$$

$$X_1 = 50$$

Q4)

$$L(\hat{w}) = -\frac{1}{n} \sum_{i=1}^n \log \frac{e^{\eta_{y_i}(\hat{x}_i)}}{\sum_{j=1}^K e^{\eta_j(\hat{x}_i)}}$$

$$= -\frac{1}{n} \sum_{i=1}^n \left[\eta_{y_i}(\hat{x}_i) - \log \left(\sum_{j=1}^K e^{\eta_j(\hat{x}_i)} \right) \right]$$

$$\frac{\partial L(\hat{w})}{\partial \eta_k(\hat{x}_i)} = -\frac{1}{n} \left[1_{\{k=y_i\}} - \frac{e^{\eta_k(\hat{x}_i)}}{\sum_{j=1}^K e^{\eta_j(\hat{x}_i)}} \right], \quad \text{where } 1_{\{k=y_i\}} = \begin{cases} 1, & \text{if } k=y_i \\ 0, & \text{otherwise} \end{cases}$$

$$= -\frac{1}{n} \left[1_{\{k=y_i\}} - \hat{y}_{ik} \right]$$

$$\eta_k(\hat{x}_i) = \hat{w}_k^T \hat{x}_i$$

$$\frac{\partial \eta_k(\hat{x}_i)}{\partial \hat{w}_k} = \hat{x}_i$$

$$\begin{aligned} \frac{\partial L(\hat{w})}{\partial \hat{w}_k} &= \sum_{i=1}^n \frac{\partial L(\hat{w})}{\partial \eta_k(\hat{x}_i)} \cdot \frac{\partial \eta_k(\hat{x}_i)}{\partial \hat{w}_k} \\ &= -\frac{1}{n} \sum_{i=1}^n [1_{\{k=y_i\}} - \hat{y}_{ik}] \hat{x}_i \end{aligned}$$

Stacking gradients from all K classes to obtain ∇L against weight matrix W :

$$\begin{aligned} \nabla_{\hat{W}} L(\hat{w}) &= \frac{1}{n} \sum_{i=1}^n [\hat{y}_i - y_i] \hat{x}_i^T \\ &= \frac{1}{n} \left[\text{softmax}(\hat{W} \hat{X}) - Y \right] X^T \end{aligned}$$

Q5)

C

$$Q6) \quad X_{yes} \sim N(10, 36)$$

$$X_{no} \sim N(0, 36)$$

$$\pi_k = 0.8, \pi_l = 0.2$$

$$\text{Let } Y = \begin{cases} 1, & \text{if company gave dividends} \\ 0, & \text{otherwise.} \end{cases}$$

$$\begin{aligned} P(Y=1 | X=14) &= \frac{\pi_k \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(X-\mu_k)^2}{2\sigma^2}}}{\sum_{k=1}^K \pi_k \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(X-\mu_k)^2}{2\sigma^2}}} \\ &= \frac{0.8 \left(\frac{1}{\sqrt{2\pi}(6)} \cdot e^{-\frac{(14-10)^2}{2(36)}} \right)}{0.8 \left(\frac{1}{\sqrt{2\pi}(6)} \cdot e^{-\frac{(14-10)^2}{2(36)}} \right) + 0.2 \left(\frac{1}{\sqrt{2\pi}(6)} \cdot e^{-\frac{(14-0)^2}{2(36)}} \right)} \\ &= 0.97989 \end{aligned}$$

Q7a) • Split X_1 at 4

$$(X_1 < 4) \quad \text{Gini}_{\text{left}} = 1 - (1^2 + 0^2) = 0$$

$$(X_1 > 4) \quad \text{Gini}_{\text{right}} = 1 - \left(\left(\frac{1}{2}\right)^2 + \left(\frac{1}{2}\right)^2 \right) = 1 - \left(\frac{1}{4} + \frac{1}{4} \right) = \frac{1}{2}$$


$$\begin{aligned} \text{Weighted Gini} &= \frac{5}{11} (0) + \frac{6}{11} \left(\frac{1}{2} \right) \\ &= \frac{3}{11} \end{aligned}$$

• Split X_1 at 6


$$(X_1 < 5) \quad \text{Gini}_{\text{left}} = 1 - (0^2 + 1^2) = 0$$

$$(X_1 > 5) \quad \text{Gini}_{\text{right}} = 1 - (1^2 + 0^2) = 0$$

$$\text{weighted Gini} = 0$$



Q7b)



Q7c) 1). The data does not look to be linearly separable based on the picture.

Therefore, to achieve 0 training error, we could attempt to use Quadratic Discriminant Analysis to fit parabolas that might separate the 2 classes with no training error.

- There are no hyperparameters for quadratic discriminant analysis.

2) Logistic regression with polynomial features for non-linear boundary. Hyperparameter to tune will be learning rate used in gradient descent.

3) K-nearest neighbour with hyperparameter $K=1$. For no training error, each training point can be classified correctly as it will always be its own nearest neighbour.

STAT3612: Statistical Machine Learning

Assignment 2: Classification

DUE: Nov 10, 2024, Sunday, 11:59 PM

```
In [1]: ! pip install numpy pandas matplotlib scikit-learn keras
```

Requirement already satisfied: numpy in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (2.0.2)

Requirement already satisfied: pandas in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (2.2.2)

Requirement already satisfied: matplotlib in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (3.5.2)

Requirement already satisfied: scikit-learn in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (1.3.0)

Requirement already satisfied: keras in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (3.8.0)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from pandas) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from pandas) (2024.1)

Requirement already satisfied: cyclor>=0.10 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (4.34.4)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (1.4.4)

Requirement already satisfied: packaging>=20.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (24.1)

Requirement already satisfied: pillow>=6.2.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (9.2.0)

Requirement already satisfied: pyparsing>=2.2.1 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from matplotlib) (3.0.9)

Requirement already satisfied: scipy>=1.5.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (1.9.0)

Requirement already satisfied: joblib>=1.1.1 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (1.3.1)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (3.2.0)

Requirement already satisfied: absl-py in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (2.1.0)

Requirement already satisfied: rich in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (13.7.1)

Requirement already satisfied: namex in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (0.0.8)

Requirement already satisfied: h5py in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (3.11.0)

Requirement already satisfied: optree in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (0.14.1)

Requirement already satisfied: ml-dtypes in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from keras) (0.4.1)

Requirement already satisfied: six>=1.5 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)

Collecting numpy

Downloading numpy-1.24.4-cp310-cp310-win_amd64.whl.metadata (5.6 kB)

Requirement already satisfied: typing-extensions>=4.5.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from optree->keras) (4.10.0)

Requirement already satisfied: markdown-it-py>=2.2.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from rich->keras) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from rich->keras) (2.15.1)

Requirement already satisfied: mdurl~>=0.1 in c:\users\kevin\appdata\local\programs\python\python310\lib\site-packages (from markdown-it-py>=2.2.0->rich->keras) (0.1.2)

Downloading numpy-1.24.4-cp310-cp310-win_amd64.whl (14.8 MB)

```

----- 0.0/14.8 MB ? eta -:--:--
----- 1.8/14.8 MB 8.4 MB/s eta 0:00:02
----- 8.1/14.8 MB 20.2 MB/s eta 0:00:01
----- 13.6/14.8 MB 22.0 MB/s eta 0:00:01

```

```
----- 14.8/14.8 MB 21.7 MB/s eta 0:00:00
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 2.0.2
    Uninstalling numpy-2.0.2:
      Successfully uninstalled numpy-2.0.2
Successfully installed numpy-1.24.4
```

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

tensorflow-intel 2.18.0 requires numpy<2.1.0,>=1.26.0, but you have numpy 1.24.4 which is incompatible.

```
In [10]: import numpy as np
from sklearn.datasets import fetch_openml, load_digits
from sklearn.model_selection import train_test_split
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import accuracy_score, precision_score, f1_score, recall_score
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.utils import shuffle
```

```
In [2]: # Q7 (a)

# -----
# Write your code here
# Load the MNIST dataset
X, y = fetch_openml('mnist_784', version=1, return_X_y=True, as_frame=False)
y = y.astype(int)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=60000, test_size=10000)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_train[0])

# Fit LDA model
lda = LinearDiscriminantAnalysis()
lda.fit(X_train, y_train)

# Make prediction and access accuracy
y_pred = lda.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Testing Data Accuracy: {accuracy * 100:.2f}%")

# -----
```

c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\datasets_openml.py:1002: FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.

```
warn(
(60000, 784)
(10000, 784)
(60000,)
3
Testing Data Accuracy: 86.21%
```

```
In [ ]: # Q7 (b)

# -----
# Write your code here
# Apply filter
y_7b = y[(y==0) | (y==1)]
X_7b = X[(y==0) | (y==1)]

X_train_7b, X_test_7b, y_train_7b, y_test_7b = train_test_split(X_7b, y_7b, test_si

# Initialize LDA
lda = LinearDiscriminantAnalysis()
lda.fit(X_train_7b, y_train_7b)

y_pred_7b = lda.predict(X_test_7b)

# Compute metrics and display results
precision = precision_score(y_test_7b, y_pred_7b)
recall = recall_score(y_test_7b, y_pred_7b)
f1 = f1_score(y_test_7b, y_pred_7b)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
# -----

Precision: 0.9950
Recall: 0.9944
F1-Score: 0.9947
Accuracy: 0.9942
```

```
In [14]: from keras.datasets import cifar10

(X_train, y_train), (X_test, y_test) = cifar10.load_data()
print('X_train shape:', X_train.shape)
print('X_test shape:', X_test.shape)
print('y_train shape:', y_train.shape)
print('y_test shape:', y_test.shape)

X_train = X_train.reshape(X_train.shape[0], -1)
X_test = X_test.reshape(X_test.shape[0], -1)

# Normalise pixel values to [0,1]
X_train = X_train / 255
X_test = X_test / 255

X_train shape: (50000, 32, 32, 3)
X_test shape: (10000, 32, 32, 3)
y_train shape: (50000, 1)
y_test shape: (10000, 1)
```

Consider flatten the images to 1D vectors and make the shape of the X be (Batch, dim).

```
In [ ]: # Q8 (a)

# -----
# Write your code here
# Flatten images to 1D vector
X_train = X_train.reshape(X_train.shape[0], -1)
X_test = X_test.reshape(X_test.shape[0], -1)

# Normalise pixel values to [0,1]
X_train = X_train / 255
X_test = X_test / 255
```



```

print(X_train)

# Standardize the data since L1 penalty is to be used.
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Define a range of regularization strengths (C values)
C_values = np.logspace(-2, 2, 5) # Logarithmic range from 1e-4 to 1e4
print(C_values)
accuracies = []

# Train and evaluate the model for each C value
for C in C_values:
    print(f"Training with C={C:.4f}...")
    model = LogisticRegression(
        multi_class="multinomial",
        solver="saga",
        penalty="l1",
        C=C, # Inverse of regularization strength in sklearn
        max_iter=100,
        random_state=42,
        warm_start=True,
    )
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    accuracies.append(accuracy)
    print(f"Accuracy: {accuracy:.4f}")

# -----

```

```

X_train shape: (50000, 3072)
X_test shape: (10000, 3072)
[[0.23137255 0.24313725 0.24705882 ... 0.48235294 0.36078431 0.28235294]
 [0.60392157 0.69411765 0.73333333 ... 0.56078431 0.52156863 0.56470588]
 [1.          1.          1.          ... 0.31372549 0.3372549 0.32941176]
 ...
 [0.1372549 0.69803922 0.92156863 ... 0.04705882 0.12156863 0.19607843]
 [0.74117647 0.82745098 0.94117647 ... 0.76470588 0.74509804 0.67058824]
 [0.89803922 0.89803922 0.9372549 ... 0.63921569 0.63921569 0.63137255]]
[1.e-02 1.e-01 1.e+00 1.e+01 1.e+02]
Training with C=0.0100...

```

```

c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().

```

```

y = column_or_1d(y, warn=True)
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

```

```

Accuracy: 0.4025

```

```

Training with C=0.1000...

```

```

c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().

```

```

y = column_or_1d(y, warn=True)
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(

```

Accuracy: 0.4124
 Training with C=1.0000...

```
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

Accuracy: 0.3923
 Training with C=10.0000...


```
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
```

Accuracy: 0.3895
 Training with C=100.0000...

```
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1184: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model\_sag.py:350: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
  warnings.warn(
```

Accuracy: 0.3898

```
In [4]: # Plot the accuracy vs. weight of the L1 penalty
plt.figure(figsize=(10, 6))
plt.plot(C_values, accuracies, marker='o', linestyle='--', color='b')
plt.xscale('log') # Use logarithmic scale for C values
plt.xlabel("Regularization Strength (C)")
plt.ylabel("Accuracy")
plt.title("Accuracy vs. Weight of L1 Penalty (C) for Softmax Regression")
plt.grid(True)
plt.show()
```



The L1 penalty term is defined as:

$$L_{L1} = \lambda \|\mathbf{W}\|_1 = \lambda \sum_{i,j} |w_{ij}|$$

The gradient of the L1 penalty with respect to the weights is:

$$\nabla_{\mathbf{W}} \mathcal{L}_{L1} = \lambda \cdot \text{sign}(\mathbf{W})$$

, where $\text{sign}(\mathbf{W})$ is the element-wise sign of the weight matrix.

```
In [54]: # Q8 (b)
np.random.seed(0)

def softmax(x):
    return np.exp(x) / np.sum(np.exp(x), axis = 1, keepdims=True)

def sgd_step(w, X_train_mini, y_train_mini, lr, minibatch_size, lambda_l1):
    """
    Perform a single step of Stochastic Gradient Descent (SGD) to update the weight
    matrix w.

    Parameters:
    w: ndarray
        Weight matrix of the model. This is the parameter we want to update.
    X_train_mini: ndarray
        A mini-batch of training data. Each row corresponds to a training example.
    y_train_mini: ndarray
        One-hot encoded labels for the mini-batch. Each row is the label for the corresponding example.
    lr: float
        Learning rate for the gradient descent step.
    minibatch_size: int
        The size of the mini-batch used in the update step.
    lambda_l1: float
        Regularization strength for the L1 penalty.

    Returns:
    w: ndarray
        Updated weight matrix after the SGD step.
    """
```

```

# Compute the gradient of the cross-entropy loss with respect to the weights
dw = (softmax((w @ X_train_mini.T).T).T - y_train_mini.T) @ X_train_mini / minibatch_size

# Add the gradient of the L1 regularization term
dw += lambda_l1 * np.sign(w)

# Update the weights by subtracting the learning rate times the gradient
w = w - lr * dw

# Return the updated weight matrix
return w

def sgd(w, X_train, y_train, lr, lambda_l1, minibatch_size=64):
    """
    Perform Stochastic Gradient Descent (SGD) over the entire dataset with L1 regularization.

    Parameters:
    w: ndarray
        Weight matrix of the model. This is the parameter we want to update.
    X_train: ndarray
        Training data. Each row corresponds to a training example, and each column corresponds to a feature.
    y_train: ndarray
        One-hot encoded labels for the training data. Each row is the label for the corresponding example.
    lr: float
        Learning rate for the gradient descent step.
    minibatch_size: int, optional
        The size of the mini-batch used in each update step. Default is 64.
    lambda_l1: float, optional
        Regularization strength for the L1 penalty. Default is 0.01.

    Returns:
    w: ndarray
        Updated weight matrix after processing the entire training data using SGD.
    """
    # Shuffle the dataset to ensure that the mini-batches are selected randomly in each epoch
    X_train, y_train = shuffle(X_train, y_train)

    # Iterate over the dataset in mini-batches
    for i in range(0, X_train.shape[0], minibatch_size):
        # Extract the current mini-batch from the training data
        X_train_mini = X_train[i:i + minibatch_size]
        # Extract the corresponding mini-batch of labels
        y_train_mini = y_train[i:i + minibatch_size]

        # Perform a single SGD step to update the weights using the mini-batch
        w = sgd_step(w, X_train_mini, y_train_mini, lr, minibatch_size, lambda_l1)

    # Return the updated weights after all mini-batches have been processed
    return w

def train_with_SGD(model, X_train, y_train, lr, epoch_num, lambda_l1):
    """
    Train the model using Stochastic Gradient Descent (SGD) with L1 regularization.

    Parameters:
    model: dict
        Dictionary containing the weight matrix.
    X_train: ndarray
        Training data. Each row corresponds to a training example, and each column corresponds to a feature.
    y_train: ndarray
        One-hot encoded labels for the training data. Each row is the label for the corresponding example.
    lr: float
        Learning rate for the gradient descent step.
    epoch_num: int
        Number of epochs to train the model.
    """

```

```

        Number of epochs to train the model.
    lambda_l1: float
        Regularization strength for the L1 penalty.

    Returns:
    model: dict
        Updated model after training.
    """
    # Extract the weight matrix from the model
    w = model['w']

    # Train the model for the specified number of epochs
    for epoch in range(epoch_num):
        w = sgd(w, X_train, y_train, lr, lambda_l1=lambda_l1)
        if (epoch + 1) % 10 == 0:
            print(f"Epoch {epoch + 1}/{epoch_num}")

    # Update the model with the trained weights
    model['w'] = w
    return model

def test(model, X_test, y_test):
    """
    Evaluate the model on the test set.

    Parameters:
    model: dict
        Dictionary containing the model parameters (e.g., weight matrix).
    X_test: ndarray
        Test data. Each row corresponds to a test example, and each column corresponds to a feature.
    y_test: ndarray
        One-hot encoded labels for the test data. Each row is the label for the corresponding example.

    Returns:
    accuracy: float
        Accuracy of the model on the test set.
    """
    # Extract the weight matrix from the model
    w = model['w']

    # Compute the predicted probabilities
    scores = np.dot(X_test, w.T)
    y_pred = np.argmax(scores, axis=1)
    y_true = np.argmax(y_test, axis=1)

    # Compute the accuracy
    accuracy = np.mean(y_pred == y_true)
    return accuracy

# -----
# Write your code here
# One-hot encode labels
encoder = OneHotEncoder(sparse=False)
y_train_onehot = encoder.fit_transform(y_train.reshape(-1, 1))
y_test_onehot = encoder.fit_transform(y_test.reshape(-1, 1))

# Initialize the model
n_features = X_train.shape[1]
n_classes = y_train.shape[1]
model = {'w': np.random.randn(n_classes, n_features) * 0.01}

```

```
# Train the model
model = train_with_SGD(model, X_train, y_train_onehot, lr=0.002, epoch_num=100, lan

# Test the model
accuracy = test(model, X_test, y_test_onehot)
print(f"Test Accuracy: {accuracy:.4f}")

# -----
```

```
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\preprocessing\_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
```

```
warnings.warn(
c:\Users\kevin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\preprocessing\_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
```

```
warnings.warn(
```

```
Epoch 10/100
```

```
Epoch 20/100
```

```
Epoch 30/100
```

```
Epoch 40/100
```

```
Epoch 50/100
```

```
Epoch 60/100
```

```
Epoch 70/100
```

```
Epoch 80/100
```

```
Epoch 90/100
```

```
Epoch 100/100
```

```
Test Accuracy: 0.4031
```

```
In [55]: # Find the class with the highest accuracy
class_accuracies = []
for k in range(n_classes):
    mask = np.argmax(y_test, axis=1) == k
    class_accuracies.append(np.mean(y_pred[mask] == k))

best_class = np.argmax(class_accuracies)
print(f"Class with highest accuracy: {best_class}")

# Visualize an image from the best class
best_class_indices = np.where(np.argmax(y_test, axis=1) == best_class)[0]
sample_index = best_class_indices[0]
sample_image = X_test[sample_index].reshape(32, 32, 3)

plt.imshow(sample_image)
plt.title(f"Class {best_class}")
plt.axis("off")
plt.show()
```

```
Class with highest accuracy: 0
```

Class 0

